



Paralegal: Practical Static Analysis for Privacy Bugs

Justus Adam, Carolyn Zech, Livia Zhu, Sreshtaa Rajesh, Nathan Harbison Mithi Jethwa, Will Crichton, Shriram Krishnamurthi, Malte Schwarzkopf

Brown University



BROWN

Problem

Existing tools are not practical enough to help developers find subtle privacy bugs.

Practical tools must meet these challenges:

1. Support for relevant privacy policies
2. No changes to code, behavior & performance
3. Policy robust against source code changes
4. Support libraries w/o needing complex models

Approach	1. Flexible	2. Non-invasive	3. Robust	4. Libraries
IFC	✗	✗	✓	✗
Specialized Bug Finder	✗	✓	✓	✓
Runtime Checks	✓	✗	✓	✓
Code Query Engine	✓	✓	✗	✗
Paralegal	✓	✓	✓	✓

Policy: expressive first-order logic policies with legal-document-like syntax

Succinct and auditable

Somewhere:

1. For each "data type" marked `user_data`:
 - A. There is a "source" that produces "data type" where:
 - a. There is a "delete query" marked `deletes`:
 - i) "source" goes to "delete query"

Markers: Semantically meaningful terms that connect policy and application code

```
#[paralegal::marker(user_data)]
struct Post { ... }
#[paralegal::marker(user_data)]
struct Comment { ... }
```

Policy does not need to change when code does

Real bug in Plume app: Devs forgot to delete comments

Link

Application Code




Application Developer

Fixes

Representation:

Program Dependence Graph with markers



Handles library modeling, scaling and FFI

Paralegal

Emits

Paralegal

```
'Deletion Policy' not satisfied.
No entrypoints satisfied rule 1.A

Entrypoint `delete_user`
Did not satisfy Rule 1.A
  There is a "source"
  that produces "data type"

for the "data type" type: `Comment` (Rule 1)
```

Approximation: Rusts types let Paralegal approximate a PDG for any function.

```
fn HashMap::get_mut<'a, 'k, K, V>(&'a mut self,
    key: &'k K) -> Option<'a mut V>
```

- `mut` tells Paralegal only `self` is modified, not `key`
- lifetimes `'a` and `'k` tell Paralegal that `self` and `key` cannot alias

Paralegal finds subtle privacy bugs in real code bases by

1. separating concerns between policy writers and developers; and
2. leveraging Rust for precise approximations

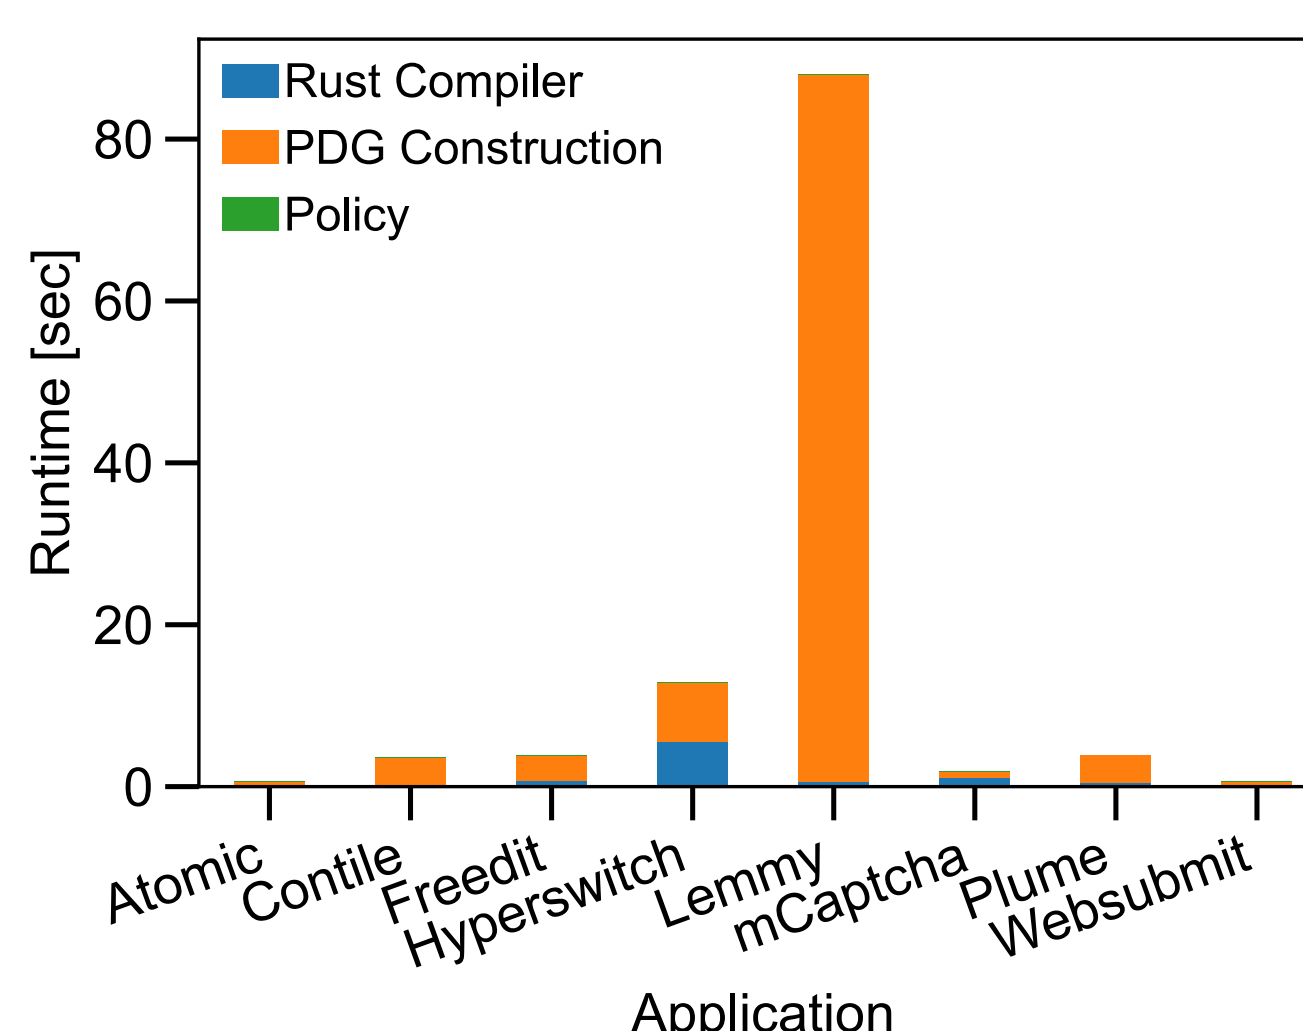
Evaluation

Setup: 8 Applications and 11 Policies

(Access Control, Purpose Limitation, Data Deletion, Credential Security, ...)

Bugs: Found 7 subtle bugs. 2 new, 5 previously known

Related Work Comparison: IFC fails to express half of the policies. CodeQL only finds half the bugs



Performance

- Full codebase check: <90 seconds
⇒ fast enough for CI
- Incremental check: <5 seconds
⇒ fast enough for IDE

Paralegal is open source.

github.com/brownsys/paralegal